

Business Requirements Document

Client Onboarding & KYC Workflow Optimization

1. Business Case

The institution currently applies a one-size-fits-all identity verification process to all onboarding applicants. This creates two simultaneous, opposing costs: (a) avoidable conversion loss from low-risk users abandoning a disproportionately heavy verification flow, and (b) avoidable regulatory exposure from not explicitly differentiating high-risk and PEP applicants for enhanced scrutiny. A risk-based tiering approach resolves both by allocating friction proportionally to actual risk rather than uniformly.

2. Current-State Pain Points

- Uniform verification flow regardless of applicant risk profile (assumption: this is the as-is state for this project's fictional institution)
- No explicit, automated PEP/sanctions-driven routing to Enhanced Due Diligence — PEP status, if checked at all, does not change the applicant's path
- Estimated drop-off concentrated at the document-verification stage (planning assumption: ~28%, see Problem Statement — to be validated against mock-dataset funnel analysis on Day 4-5, not treated as a real measurement)
- No quantified Time-to-Onboard (TTO) differentiation by risk tier, making it impossible to know whether friction is correctly allocated

3. Target-State Goals

- Applicants are scored at intake using a transparent, deterministic weighted formula (Country of Origin, Transaction Velocity, Verification Match Rate)
- Applicants are tiered: **Low, Medium, High** — with PEP status automatically forcing **High / EDD** regardless of computed score
- Low-risk applicants experience materially shorter TTO than the current uniform flow
- High-risk and PEP applicants are routed to Enhanced Due Diligence with full audit traceability
- The scoring threshold separating tiers is chosen based on actual analysis of score distribution and outcome data — not an arbitrary number

4. Success Metrics

Metric	Target	Source
TTO — Low-risk tier	Materially faster than current uniform baseline (exact figure set after Day 4-5 mock-data analysis)	SQL query, Day 5
TTO — High-risk/EDD tier	Slower than Low tier — expected and acceptable, since this is deliberate scrutiny, not failure	SQL query, Day 5
Drop-off reduction	Directional target only (no real measurement possible without production data)	Stated assumption
Verification Pass Rate by tier	Reported as-observed from mock dataset, no pre-set target	SQL query, Day 5
Threshold justification	Backed by real false-positive/false-negative counts at candidate thresholds (e.g., 70 vs 75)	Python script, Day 4

5. Assumptions & Constraints

- This is a portfolio/demonstration project; no real production data, users, or vendor integration exists
- The ~28% drop-off figure is a planning assumption for sizing the business case, not a measured fact, and is explicitly labeled as such everywhere it appears
- The scoring model is a transparent weighted formula, not a machine-learning model, by design — chosen for explainability to Compliance and regulators
- Exact weights and the approval threshold are placeholders until the Day 4 threshold-tuning exercise produces real output; the FRD will be patched, not rewritten, once those numbers exist
- No real third-party KYC vendor will be contacted or integrated; the API spec (Day 5) is a hand-off artifact only

6. Out of Scope

- No OCR, biometric, or facial-matching technology implementation
- No actual third-party API call or live integration
- No machine learning of any kind