

Problem Statement

Our institution currently reconciles transactions manually across an internal ledger, a payment gateway feed, and an external clearing-house settlement file. Operations staff pull CSV exports from all three systems and use Excel-based lookups to identify mismatches — a process that is slow, inconsistent between operators, and produces no standardized audit trail of how or why a given mismatch was resolved.

Stated assumption (flagged, not a finding): for the purposes of sizing this project's business case, we assume operations staff spend an average of ~3.5 hours per day on manual reconciliation, with an estimated 15% of exceptions remaining unresolved past 72 hours at any given time. These figures are planning assumptions used to size the BRD below — they are not measured facts, since no real production reconciliation data exists for this portfolio project, and no script or query in this project will be run to validate or revise them. Unlike Project A, this project has no later stage where real numbers would supersede these assumptions; they remain stated assumptions throughout the entire deliverable.

The core ask: Design a rules-based matching and exception-handling framework that automatically resolves transactions which match within a clearly defined, low-risk tolerance, while routing genuine exceptions into an aging-based risk-bucket system (0–24h / 24–72h / >72h) with a mandatory audit trail for every resolution — whether automated or manual. This resolves the Finance-vs-Compliance tension by changing *what* gets fast-tracked, not by making the whole process faster or slower uniformly.